
Image Classification with Deep Learning

Timoshenko Dmitrii
Department of Statistics
University of California, Berkeley
timoshenko_dmitrii@berkeley.edu

Abstract

This project explores various Deep Learning approaches for image classification in sports. The dataset includes 100 classes of different sports, with a limited number of examples for training, validation, and testing. We experiment with different CNN architectures, including state-of-the-art models, to evaluate their performance on the test set. Additionally, we investigate techniques to improve model accuracy, such as data augmentation, increasing the depth of classifier layers, and blending with boosting. Through our experiments, we aim to identify the best-performing methods for sports image classification.

1 Introduction

Image classification is an important task in computer vision, with numerous applications in fields such as robotics, surveillance, and medical imaging. In recent years, deep learning techniques have been at the forefront of image classification, achieving state-of-the-art results in a variety of domains. Sports imagery is a particularly interesting domain for image classification, as it poses unique challenges due to the high variability of the images and the need to distinguish between similar poses and actions. While there have been several studies on sports image classification, this paper aims to explore different methods for this task. Specifically, the performance of various deep learning architectures and preprocessing techniques will be investigated on a dataset of images from different sports. The goal is to provide insights into which methods are most effective for sports image classification and to contribute to the development of more accurate and robust classification systems for sports imagery.

2 Data

The key aspect of this project was collecting a suitable dataset to train the models. The dataset of 100 different sports was obtained from Kaggle¹, which consists of 100 training samples, 5 validation samples, and 5 test samples for each class. The images in the dataset are colorful and non-centered, with the color scheme being RGB. Each image comprises three channels of red, green, and blue colors, which are combined to create the final image. However, the non-centered nature of the data suggests that the images may not be perfectly aligned and may require some preprocessing before they can be used in machine-learning models. Moreover, unlike datasets like MNIST, where the images are well-centered, some traditional computer vision methods may not show good results in sports image classification tasks. Therefore, I used deep learning approaches, which I discuss in detail later in the paper.

To provide a better understanding of the dataset, Figure 1 shows some examples of the images.

¹<https://www.kaggle.com/datasets/gpiosenka/sports-classification>



Figure 1: Examples of images in the dataset

3 Methods

In this section, I describe the methods I used to achieve the best performance. There are some reasons, why traditional machine learning methods can't be used:

- The data is non-centered, making it difficult for traditional methods such as tree-based and regression models to create appropriate boundaries to classify the images effectively
- The images in the dataset have a high resolution of 224x224x3, making it computationally expensive to use traditional models even after grayscaleing.
- Color is a critical feature in sports image classification. For instance, the color difference between soccer and basketball images is significant and plays a crucial role in improving the prediction power of the models.

Given the challenges posed by the sports image classification task, I decided to use deep learning methods, and specifically, a Convolutional Neural Network (CNN), as it is one of the most suitable models for this task.

3.1 Pretrained CNN: ConvNet as fixed feature extractor

Motivation. Since the dataset is relatively small, building a neural network from scratch would not be an ideal approach. Deep learning methods typically require a large amount of data to achieve good performance.² Therefore, I decided to use pre-trained CNN models instead. The main idea behind this approach is to use a pre-trained CNN model with frozen main layers (known as CNN codes), and only train the classifier layer with randomly initialized weights.

Training. I used various CNN architectures as fixed feature extractors and modified only the number of neurons in the last layer to correspond with the total number of classes in the dataset (i.e., 100). Each model was trained on a dataset consisting of 100 samples for each class and validated on 5 images for each class after every epoch to evaluate the model's performance and prevent overfitting. Prior to training, the images were normalized to enhance model performance.

Main parameters of the training procedure described in Table 1. For calculation a loss after each epoch, I used cross entropy (Eq. 1), with averaging it over number of samples in batch.

$$Loss_n = -w_{y_n} \log \frac{\exp(x_{n,y_n})}{\sum_{c=1}^C \exp(x_{n,c})} \quad (1)$$

I chose different models, including the most recent one, to check their performance on my dataset. The training was completed on an NVIDIA A100 GPU, and the training time in Table 2 is measured in minutes. The accuracy metric for the test is defined as the share of correctly predicted images in the test dataset (Eq. 2).

I chose several pretrained models to test their performance. Before discussing the results, I will briefly explain their architectures.

²However, I trained my own model. Accuracy was too low: 0.296, which is higher than I get from ShuffleNet

Table 1: Training parameters for CNN models

Parameter	Value	Comment
Batch size	32	Number of images to proceed in one forward step
No. of epoches	100	Total number of epoches. All models were fully trained before the number
Early stop	5	If no improvement on validation dataset for this number of epochs, training will be stopped
Loss function	Cross entropy	Most common criteria in classification
Optimization algorithm	SGD/Adam	Optimizer used in the model. For some models SGD worked better than Adam
Learning rate	0.001	Regulates the speed of the weight adjustment
Weights	ImageNet	All weights for pre-trained models were frozen. Last layer had random weight initialization

$$Accuracy = \frac{\sum_{i=1}^n \mathbb{1}(\hat{y}_i \neq y_i)}{n} \quad (2)$$

3.1.1 AlexNet

AlexNet is a deep convolutional neural network developed by Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton [3]. It consists of 8 layers, including 5 convolutional layers and 3 fully connected layers. The input to the network is an image of size 224x224x3, where 3 represents the RGB channels. This model serves as a good baseline, as it learns quickly and provides an expected accuracy threshold.

3.1.2 ResNet

ResNet (short for "Residual Network") is a deep convolutional neural network architecture that was introduced by Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun in 2015 [2]. It was designed to address the problem of vanishing gradients in very deep neural networks by introducing skip connections that allow information to bypass certain layers in the network. In many applications, this model shows great results with relatively low computational costs.

3.1.3 Inception

The Inception architecture, also known as GoogleNet, is a convolutional neural network architecture that was introduced by researchers at Google in 2014 [5]. The authors proposed some key differences from previous networks:

- Increase the model not only in depth, but also in width
- Reduce the number of layers with a transition from a large number of parameters to a smaller number, i.e. avoid bottlenecks
- Small convolutions can be effective as nearby pixels are often highly correlated, so reducing the image size can be done without significant losses in accuracy

The architecture differs greatly from AlexNet and ResNet, so I decided to test its performance on my dataset.

3.1.4 ShuffleNet

ShuffleNet is a convolutional neural network architecture that was introduced by researchers at Megvii (also known as Face++) in 2017 [4]. ShuffleNet uses a channel shuffle technique to reduce

Table 2: Training Results

Model	Metrics	
	Accuracy on test	Training Time
AlexNet (2012)	0.86	8
ResNet50 (2015)	0.97	23
ResNet152 (2015)	0.94	38
Inception V3 (2015)	0.92	47
ShuffleNet V2 (2018)	0.03	21
EfficientNet B3 (2019)	0.88	17
EfficientNet B6 (2019)	0.80	24
EfficientNetV2-S (2021)	0.87	36
EfficientNetV2-L (2021)	0.94	34
ViT-B (2021)	0.95	24
ViT-L (2021)	0.92	89

the number of parameters and the computational cost of the network while balancing accuracy with computational efficiency.

3.1.5 EfficientNet

EfficientNet is a convolutional neural network architecture that was introduced by researchers at Google in 2020 [6, 7]. The EfficientNet architecture achieves this goal by using a compound scaling approach, where the network is scaled up or down in terms of depth, width, and resolution based on a set of scaling coefficients. The model demonstrated good performance in similar problems.

3.1.6 Vision Transformer (ViT)

The Vision Transformer (ViT) is a deep learning architecture introduced by Google in 2020 [1]. It is based on the Transformer architecture that was originally proposed for natural language processing tasks, but has been adapted for use in computer vision tasks. The ViT architecture works by transforming an input image into a sequence of tokens that are then processed by a series of Transformer blocks. Novel architecture that I also wanted to test on my dataset.

Results. Table 2 summarizes the training procedure and shows that the ResNet50 architecture performed the best among all models tested. Despite parameter tuning attempts using different optimizers, learning rates, and schedulers, the ShuffleNet V2 model did not improve its predictive power. Furthermore, larger versions of some models (ResNet, EfficientNet, ViT) did not outperform their smaller counterparts as expected. One possible reason for this is that larger models are better suited for recognizing more complex patterns in images, which may not be necessary for this particular dataset.

Also, Figure 2 demonstrates the wrongly labeled images. Overall, the results seem very good—even for humans, it is possible to make a mistake on these images.

3.2 Pretrained CNN: Data Augmentation

Motivation. Since the initial dataset is small, increasing the number of samples with augmentation is a reasonable way to improve the predictive power of CNN models.

Training. Two versions of augmentation were implemented: one using two filters, each with a probability of 0.5, and the other using several filters with different probabilities.³ The description of filters can be found in Table A1, and Figure 3 shows the filters used in training. Some filters were used with a probability less than 0.5 because they could change the picture too much for the models, so they were included with low probability to prevent overfitting. Filters that could change colors too

³Augmentation was used only on the training dataset.



Figure 2: Examples of wrong predictions

much were not used because color is one of the most important features of the dataset, as mentioned in the motivation for using pretrained CNN models.

Results. The most important outcomes of the method are as follows:

- The training time increased for most models, as expected, because transforming images requires more time to process them before actual training.
- For some models (Inception V3, EfficientNetV2-S/L, AlexNet), the quality of predictions decreased. Again, the reason for this is the small size of the initial dataset.
- Unexpectedly, the training time for some models with several filters increased compared to two filters. This was likely because I wasn't the only user of the GPU and couldn't control its usage.
- Overall, in most cases, using two filters for the given dataset is a better option, as models are less likely to miss details and don't require much time to train. The latter could be improved with the use of other technical solutions, such as Albumentations (example: <https://albumentations.ai>).



Figure 3: Examples of augmentation

3.3 ResNet50: Different Classifier

Motivation. Another method I used to attempt to increase the predictive power was to increase the depth of the classifier, as the feature extraction might be good enough to improve it further.

Training. I used the ResNet50 architecture as a feature extractor since it showed the best accuracy among all models, and its training time is not as long as that of larger models. To make the classifier layer deeper, I included more fully connected and dropout layers with ReLU and ELU as non-linear activation functions.⁴ The loss function and most of the parameters are the same as in previous sections. The changes to the initial architecture are illustrated in Figure 4 (visuals created by the author).

Results. Unfortunately, making the classifier deeper did not improve the predictive power as expected. In all cases, accuracy decreased by 1.5-2 percentage points compared to the initial ResNet50 architecture as a feature extractor. I tried different specifications and have described them in Table A3.

⁴Some clarifications on the notation: D - dropout; L - FC layer; Re - ReLU activation function; Elu - ELU activation function




Figure 4: ResNet50 as a feature extractor, combined with newly trained classifier

3.4 Blending: CNN with Boosting

Motivation. Another method I attempted to improve the classification power of the models was to combine their predictions to build one strong classifier. I used trained models as a feature extractor and built another model that identifies the most important features and uses them to make a final decision on the class label. Figure 5 demonstrates this idea.

Training. The training procedure is as follows:

1. Train five different neural networks.
2. Instead of predicting the class, obtain the outputs before the softmax layer. This yields 100 features from each model for each image, resulting in 500 features for each picture.
3. Without losing the train/valid/test split, train a boosting model to increase the overall accuracy of the models.
4. Predict the class of the test images and evaluate the new accuracy.

I tried different specifications and I describe them in Table A4.

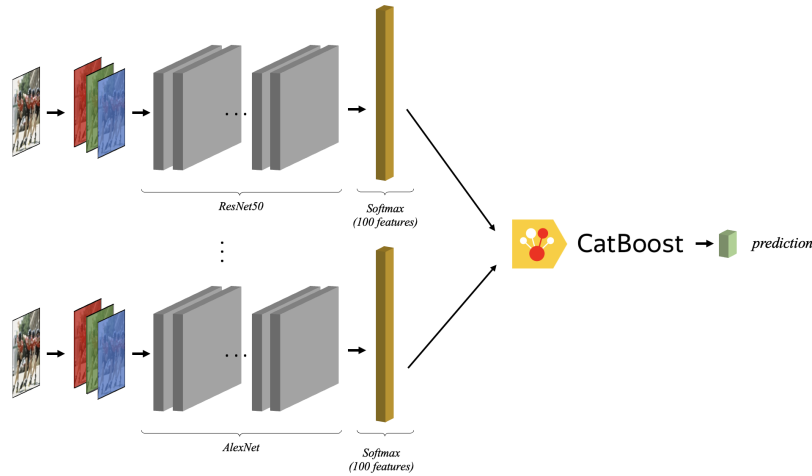


Figure 5: Several CNN models combined to make a prediction after boosting

Results. This method didn't increase overall accuracy of models, even given the fact that in data there was an output from ResNet50 architecture, which showed 0.97 accuracy. Interesting, that after checking confusion matrix, I noticed the model predict `air_hockey` too often: among 51 errors for different images, the model wrongly predicted `air_hockey`, instead of correct class. Possible solution here: drop this class, but even after removing the class, nothing changes – accuracy is still relatively low, 0.84.

4 Conclusion

In this project, I tried different models for image classification and found that the ResNet50 architecture performed very well on my data, achieving high performance (0.97) in a relatively short training time. Various techniques that are usually employed to improve model accuracy didn't yield significant improvements, possibly due to the small size of the dataset and the already high accuracy of the ResNet50 model. Areas for potential future research could include:

- Parameter tuning for models.
- Unfreeze layers and train them as well (now it is impossible due to high computational costs). A good thing to try would be different learning rates for layers – higher for the classifier and lower for the feature extractor.
- Increase the dataset with other images using web parsing.

References

- [1] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., and Houlsby, N. (2021). An image is worth 16x16 words: Transformers for image recognition at scale.
- [2] He, K., Zhang, X., Ren, S., and Sun, J. (2015). Deep residual learning for image recognition.
- [3] Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. In Pereira, F., Burges, C., Bottou, L., and Weinberger, K., editors, *Advances in Neural Information Processing Systems*, volume 25. Curran Associates, Inc.
- [4] Ma, N., Zhang, X., Zheng, H.-T., and Sun, J. (2018). Shufflenet v2: Practical guidelines for efficient cnn architecture design.
- [5] Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., and Wojna, Z. (2015). Rethinking the inception architecture for computer vision.
- [6] Tan, M. and Le, Q. V. (2020). Efficientnet: Rethinking model scaling for convolutional neural networks.
- [7] Tan, M. and Le, Q. V. (2021). Efficientnetv2: Smaller models and faster training.

Appendix

Table A1: Filters for augmentation

Filter	Probability	Comment	Setting 1	Setting 2
RandomHorizontalFlip	0.5	Flip an image horizontally	Yes	Yes
RandomVerticalFlip	0.5	Flip an image vertically	Yes	Yes
RandomEqualize	0.1	Equalizes the histogram of image	No	Yes
RandomPerspective	0.3	Performs perspective transform	No	Yes
RandomAutocontrast	0.3	Applies autocontrast	No	Yes

Table A2: Training Results

Model	Native		Two filters		Multiple filters	
	Training time	Accuracy	Training time	Accuracy	Training time	Accuracy
AlexNet (2012)	9	0.87	25	0.86	9	0.80
ResNet50 (2015)	23	0.97	61	0.96	42	0.97
ResNet152 (2015)	38	0.94	114	0.93	88	0.94
Inception V3 (2015)	47	0.92	163	0.78	110	0.79
SuffleNet V2 (2018)	21	0.03	112	0.03	139	0.02
EfficientNet B3 (2019)	17	0.89	47	0.89	37	0.90
EfficientNet B6 (2019)	23	0.80	110	0.80	109	0.79
EfficientNetV2-S (2021)	36	0.88	84	0.84	63	0.86
EfficientNetV2-L (2021)	33	0.94	107	0.93	106	0.94
ViT-B (2021)	24	0.95	98	0.93	86	0.93
ViT-L (2021)	89	0.92	212	0.88	160	0.89

Table A3: Training Results

Model	DLReDL	DLReDLReDL	DLReDL	DLEluDL	DLReDLReDL
Learning Rate	0.0005	0.0001	0.001	0.0005	0.0005
Optimizer	Adam	Adam	SGD	Adam	
CNN	ResNet50	ResNet50	ResNet50	ResNet50	ResNet50
Layer 1	Dropout	Dropout	Dropout	Dropout	Dropout
Layer 2	FC	FC	FC	FC	FC
	2048→1024	2048→1024	2048→1024	2048→1024	2048→1500
Layer 3	ReLU	ReLU	ReLU	ELU	ReLU
Layer 4	Dropout	Dropout	Dropout	Dropout	Dropout
Layer 5	FC	FC	FC	FC	FC
	1024→100	1024→512	1024→100	1024→100	1500→750
Layer 6	...	ReLU	ReLU
Layer 7	...	Dropout	Dropout
Layer 8	...	FC 512→100	FC 750→100
Epochs to train	19	16	48	11	10
Accuracy	95.8	95.8	96.21	96.01	96.41

Table A4: Training results from blending

After parameter turning. Unrestricted model (when maximum depth isn't defined) didn't show good result

	Model 1	Model 2	Model 3	Model 4	Model 5
Loss function	CrossEntropy	CrossEntropy	CrossEntropy	CrossEntropy	CrossEntropy
Learning Rate	0.8	0.8	0.8	0.8	0.8
Max Depth	5	5	2	2	5
L2 Regularization	5	10	10	15	15
Accuracy	0.88	0.84	0.81	0.83	0.86